15-441: Computer Networks
Project 3: Video CDN

Lead TAs:
Alex Bainbridge <abainbri@andrew.cmu.edu>
Aneek Mukherjee <aneekm@andrew.cmu.edu>

1 Introduction

In this project you will explore how video content distribution networks (CDNs) work. In particular, you will implement a video distribution systems that uses adaptive bitrate selection (described in the *Video Distribution* lecture), DNS redirect (described in the *Content Delivery and Peer-to-Peer* lecture), and your DNS server will use OSPF to decide which server is closest to a given client using OSPF (described in the *Intra-domain Routing* lecture).

Figure 1 depicts (at a high level) what this system looks like in the real world. Clients trying to stream a video first issue a DNS query to resolve the service's domain name to an IP address for one of the CDN's content servers. The CDN's authoritative DNS server selects the "best" content server for each particular client using two algorithms: (1) a round-robin algorithm that simply distributes the equally across the servers, and (2) an algorithm that picks the server that is closest to the client measured in network hops.

Once the client has the IP address for one of the content servers, it begins requesting chunks of the video the user requested. The video is encoded at multiple bitrates; as the client player receives video data, it calculates the throughput of the transfer and requests the highest bitrate the connection can support.

Implementing an entire CDN is clearly a tall order, so let's simplify things. First, your entire system will run on one host; we are providing a network simulator (described in §5) that will allow you to run several processes with arbitrary IP addresses on one machine. Our simulator also allows you to assign arbitrary link characteristics (bandwidth and latency) to each pair of "endhosts" (processes). For this project, you will do your development and testing using a virtual machine we provide (§5).

The grade distribution and deadlines for the two checkpoints are shown below:

CP	Goal	P3 Grade	Deadline
1	Proxy and Bit rate adaptation.	60%	Nov 22, 2019
2	DNS Redirect.	40%	Dec 6, 2019

The process of team formation is the same as project 2. If you have any questions about it, please let us know ASAP. For this project, we won't use autolab for grading. We will give you test scripts for each checkpoint, and autolab is only for submission purpose. Grading will be done manually by TAs. We will have some hidden tests (like what we do in project 1 and 2), so please make sure to make your codes robust.

2 Your Video Distribution System

Figure 2 shows the pieces of the system; those shaded in gray will be written by you.

Browser. You will use an off-the-shelf web browser to play videos served by your CDN (via your proxy).

Proxy. Rather than modify the video player itself, you will implement adaptive bit rate selection in an HTTP proxy. The player requests chunks with standard HTTP GET requests; your proxy will intercept

Figure 1: In the real world...


Figure 2: Your system.

Figure 3: Your proxy should listen for browser connections on `INADDR_ANY` on the port specified on the command line. It should then connect to web servers on sockets that have been bound to the proxy's fake IP address (also specified on the command line).

these and modify them to retrieve whichever bitrate your algorithm deems appropriate. To simulate multiple clients, you will launch multiple instances of your proxy. More detail in §2.1.

Web Server. Video content will be served from an off-the-shelf web server (Apache). As with the proxy, you will run multiple instances of Apache on different fake IP addresses to simulate a CDN with several content servers. More detail in §5.

DNS Server. You will implement a simple DNS (supporting only a small portion of actual DNS functionality). Your server will respond to each request with the "best" server for that particular client. More detail in §2.3.

This project has two checkpoints. For the first checkpoint you have to implement a proxy that implements video bit rate adaptation, and in the second checkpoint you will add DNS redirect to pick the "best" server for each client, considering both server load and network distance. This section describes the features of these three components in detail. Sections 4 and 4 lists what you need to complete and hand in, and summarizes grading information for each checkpoint. Section 5 provides development information for the project.

2.1 Implementating a Proxy

You are implementing a simple HTTP proxy. It accepts connections from one web browser, modifies video chunk requests as described below, resolves the web server's DNS name (part of the HTTP request), opens a connection with the resulting server IP address, and forwards the modified request to the server. Any data (the video chunks) returned by the server should be forwarded, unmodified, to the browser.

Your proxy should listen for connections from a browser on any IP address on the port specified as a command line argument (see below). When it connects to a server, it should first bind the socket to the fake IP address specified on the command line (note that this is somewhat atypical: you do not ordinarily `bind()` a client socket before connecting). Figure 3 depicts this.

Your proxy should accept multiple concurrent connections using `select()`, as in project 1. The reason is that the browser may open multiple TCP connections to download objects (chunks in our case) in parallel.

Similarly, your proxy should open multiple connections to the server (Figure 3). **It is fine to re-use select() and HTTP parsing code from project 1.**

2.2 Video Bitrate Adaptation

Many video players monitor how quickly they receive data from the server and use this throughput value to request better or lower quality encodings of the video, aiming to stream the highest quality encoding that the connection can handle. Rather than modifying an existing video client to perform bitrate adaptation, you will implement this functionality in the HTTP proxy you implemented.

Video bit rate adaptation is described in pages 7-8 of the handout of the Video Distribution Lecture. It involves a number of steps: (1) throughput estimation, (2) choosing a bit rate,

2.2.1 Throughput Estimation

The goal of bit rate adaptation is to pick the best bit rate for each video chunk, as network conditions change. In the real world, this means that the video player picks a new rate for each chunk; this function is delegated to the proxy in our case. The proxy determines the bit rate for each chunk based on the past history of the throughputs observed for downloads of earlier chunks from the same server, i.e., the throughput estimation is for the path between your proxy and a specific server IP address. To measure the throughput for each chunk you will use a "black box" approach: you simply measure the throughput for each chunk as described below, while ignoring external information (e.g., how many videos are being streamed, how many chunks are being fetched in parallel, etc.).

Your proxy should estimate the throughput for each video chunk as follows. Note the start time, t_s, of each chunk request (i.e., include `time.h` and save a timestamp using `timeofday()` when your proxy receives a chunk request from the player). Save another timestamp, t_f, when you have finished receiving the chunk from the server. Now, given the size of the chunk, B, you can compute the throughput for the chunk, T, your proxy saw for this chunk:

$$T = \frac{B}{t_f - t_s}$$

To smooth your throughput estimation, your proxy should use an exponentially-weighted moving average (EWMA). Every time you make a new measurement (T_{new}), update your current throughput estimate as follows:

$$T_{current} = \alpha T_{new} + (1 - \alpha)T_{current} \tag{1}$$

The constant $0 \leq \alpha \leq 1$ controls the tradeoff between a smooth throughput estimate (α closer to 0) and one that reacts quickly to changes (α closer to 1). You will control α via a command line argument.

Since the bit rate depends to the path to the server, you need to maintain a separate throughput estimate for each server. When a new server is used, e.g., for a new stream, set $T_{current}$ to the lowest available bitrate for that video and server. Note that however that for CP 1, you can assume that (1) there is only one client per proxy, (2) each client client will only download one video at the time, and (3) all the chunks for a video will be fetched from the same server . That simplifies the data structures you need to maintain in the proxy. However, for CP 2, you may talk to multiple servers, so separate estimates are needed.

2.2.2 Choosing a Bitrate

Once your proxy has calculated the connection's current throughput, it should select the highest offered bitrate the connection can support. For this project, we say a connection can support a bitrate if the average throughput is at least 1.5 times the bitrate. For example, before your proxy should request chunks encoded at 1000 Kbps, its current throughput estimate should be at least 1.5 Mbps.

Your proxy should learn which bitrates are available for a given video by parsing the manifest file (the ".f4m" initially requested at the beginning of the stream). The manifest is encoded in XML; each encoding of the video is described by a <media> element, whose `bitrate` attribute you should find.

Your proxy replaces each chunk request with a request for the same chunk at the selected bitrate (in Kbps) by modifying the HTTP request's Request-URI. Video chunk URIs are structured as follows:

$$\texttt{/path/to/video/<bitrate>Seq<num>-Frag<num>}$$

For example, suppose the player requests a chunk that corresponds to fragment 3 of sequence 2 of the video Big Buck Bunny at 500 Kbps:

$$\texttt{/path/to/video/500Seg2-Frag3}$$

To switch to a higher bitrate, e.g., 1000 Kbps, the proxy should modify the URI like this:

$$\texttt{/path/to/video/1000Seg2-Frag3}$$

IMPORTANT: When the video player requests `big_buck_bunny.f4m`, your proxy should instead return `big_buck_bunny_nolist.f4m`. This file does not list the available bitrates, preventing the video player from attempting its own bitrate adaptation. You proxy should, however, fetch `big_buck_bunny.f4m` for itself (i.e., don't return it to the client) so you can parse the list of available encodings as described above.

2.2.3 Implementation Details

Logging: We require that your proxy create a log of its activity in a very particular format. After each request, it should append the following line to the log:

$$\texttt{<time> <duration> <tput> <avg-tput> <bitrate> <server-ip> <chunkname>}$$

`time` The current time in seconds since the epoch.

`duration` A floating point number representing the number of seconds it took to download this chunk from the server to the proxy.

`tput` The throughput you measured for the current chunk in Kbps.

`avg-tput` Your current EWMA throughput estimate in Kbps.

`bitrate` The bitrate your proxy requested for this chunk in Kbps.

`server-ip` The IP address of the server to which the proxy forwarded this request.

`chunkname` The name of the file your proxy requested from the server (that is, the modified file name in the modified HTTP GET message).

Running the proxy: By running `make` in the root of your submission directory, we should be able to create an executable called `proxy`, which should be invoked as follows, *even if not all arguments are functional at the first checkpoint*:

$$\texttt{./proxy <log> <alpha> <listen-port> <fake-ip> <dns-ip> <dns-port> [<www-ip>]}$$

`log` The file path to which you should log the messages described in §2.2.3.

`alpha` A float in the range [0, 1]. Uses this as the coefficient in your EWMA throughput estimate (Equation 1).

`listen-port` The TCP port your proxy should listen on for accepting connections from your browser.

`fake-ip` Your proxy should bind to this IP address *for outbound connections to the web servers*. You should *not* bind your browser listen socket to this IP address — bind that socket to `INADDR_ANY`.

`dns-ip` IP address of the DNS server.

`dns-port` UDP port DNS server listens on.

`www-ip` Your proxy should accept an optional argument specifying the IP address of the web server from which it should request video chunks. If this argument is not present, your proxy should obtain the web server's IP address by querying your DNS server for the name `video.cs.cmu.edu`.

To play a video through your proxy, point a browser on your VM to the URL `http://localhost:<listen-port>/index.html`. (You can also configure VirtualBox's port forwarding to send traffic from `<listen-port>` on the host machine to `<listen-port>` on your VM; this way you can play the video from your own web browser.)

2.3 DNS Load Balancing

To spread the load of serving videos among a group of servers, most CDNs perform some kind of load balancing. A common technique is to configure the CDN's authoritative DNS server to resolve a single domain name to one out of a set of IP addresses belonging to replicated content servers. The DNS server can use various strategies to spread the load, e.g., round-robin, shortest shortest path, or current server load (which requires servers to periodically report their statuses to the DNS server).

You will write a simple DNS server that implements load balancing two different ways: round-robin and shortest path.

2.3.1 Load Balancing Algorithms

You must implement the two load balancing strategies described below. In order for you proxy to be able to query your DNS server, you must also write an accompanying DNS resolution library (see §2.3.2 for details).

(1) Round robin. First, implement a simple round-robin based DNS load balancer. Your DNS process should take as input a list of video server IP addresses (the topology's `.servers` file — §5.2) on the command line (§2.3.2); it responds to each request to resolve the name `video.cs.cmu.edu`by returning the next IP address in the list, cycling back to the beginning when the list is exhausted.

(2) Shortest path. Next you will make your DNS server somewhat more sophisticated — must return the "closest" video server to the client based on the client's IP address. In the real world, this would be done by querying a database mapping IP prefixes to geographic locations. In your implementation, you will use shortest as an alternative. You will pretend that a link state routing protocol (e.g., OSPF) is used Internet-wide and that your DNS server participates in it. Your DNS server must process link state advertisements (LSAs), build a graph of the entire network, and run Dijkstra's shortest path algorithm on the graph to determine the closest video server for a given client.

You do not need to implement LSA flooding. We will provide you with a file containing a list of LSAs which your server would have received had you actually implemented LSA flooding. The file contains one LSA per line, formatted as follows:

<center><sender> <sequence number> <neighbors></center>

sender The IP address of the node that originated this LSA.

sequence number An integer allowing you to order the LSAs from a given sender. Each node sends multiple LSAs; you should accept only the most recent (*even if they arrive out of order*). You may assume sequence numbers don't wrap back to zero within the LSA file and that nodes in the network never reboot and start over at zero.

neighbors A comma-separated string of IP addresses denoting the sender's immediate (directly connected) neighbors.

2.3.2 Implementation Details

Your DNS implementation will consist of three components: the DNS protocol, a DNS server returning the IP address of the "best" server and a client-side resolution library.

DNS protocol: The message formats for the DNS protocol are defined in section 4.1 of RFC 1035. To make your life easier:

AA Set this to 0 in requests, 1 in responses.

RD Set this to 0 in all messages.

RA Set this to 0 in all messages.

Z Set this to 0 in all messages.

NSCOUNT Set this to 0 in all messages.

ARCOUNT Set this to 0 in all messages.

QTYPE Set this to 1 in all requests (asking for an A record).

QCLASS Set this to 1 in all requests (asking for an IP address).

TYPE Set this to 1 in all responses (returning an A record).

CLASS Set this to 1 in all responses (returning an IP address).

TTL Set this to 0 in all responses (no caching).

DNS Server: Your DNS server will operate over UDP. It will bind to an IP address and port specified as command line arguments. It need only respond to requests for `video.cs.cmu.edu`; any other requests should generate a response with `RCODE` 3. By running `make` in the root of your submission directory, we should be able to create an executable called `nameserver`, which should be invoked as follows *even if not all arguments are functional*:

> ./nameserver [-r] <log> <ip> <port> <servers> <LSAs>

-r If present, this flag indicates the server should perform round-robin load balancing instead of processing LSAs and returing the client's closest server.

log The file path to which you should log the messages as described below.

ip The IP address on which your server should listen.

port The UDP port on which your server should listen.

servers A text file containing a list of IP addresses, one per line, belonging to content servers.

LSAs A text file containing a list of LSAs, one per line, in the format described above.

Logging Like your proxy, your DNS server must log its activity in a specific format. For each valid DNS query it services, it should append the following line to the log:

> <time> <client-ip> <query-name> <response-ip>

time The current time in seconds since the epoch.

client-ip The IP address of the client who sent the query.

query-name The hostname the client is trying to resolve.

response-ip The IP address you return in response.

Resolution library: This library is used by your proxy to map the host names into IP addresses. It offers one function: `resolve()`. We have provided the interface in `mydns.h`; you are to write the accompanying implementation in a file named `mydns.c`. Your proxy will use your resolver by including `mydns.h` and calling `resolve()` at the beginning of each new connection.

3 Checkpoint 1: Bit Rate Adaptation

3.1 Requirements

You MUST complete the following:

(1) Proxy. Implement a proxy that forwards requests received from the browser to the server, and forward to chunk it receives to the browser. See §2.1 for details.

(2) Bitrate adaptation. Implement bit rate adaptation and logging as descibed in Section 2.2.

(3) Explore the behavior of your proxy. Once your proxy is working, launch two instances of it on the "dumbbell" topology (topo1) we provide. Running the dumbbell topology will also create two servers (listening on the IP addresses in `topo1.servers`); you should direct one proxy to each server. Now:

1. Start playing the video through each proxy.

2. Run the topo1 events file and direct `netsim.py` to generate a log file: `./netsim.py -l <log-file> ../topos/topo1 run`

3. After 1 minute, stop video playback and kill the proxies.

4. Gather the netsim log file and the log files from your proxy and use them to generate plots for link utilization, fairness, and smoothness. Use our `grapher.py` script to do this: `./grapher.py <netsim-log> <proxy-1-log> <proxy-2-log>`

Repeat these steps for $\alpha = 0.1$, $\alpha = 0.5$, $\alpha = 0.9$ (see §2.2.1). Compile your 9 plots, labelled clearly, into a single PDF named `writeup.pdf`, along with a brief (1-3 paragraphs) discussion of the tradeoffs you make as you vary α. We are not looking for a thorough, extensive study, just general observations. For checkpoint 1 we are only giving completion points for your writeup, so it's okay if it's still a bit preliminary.

3.2 Hand In

You will submit your code as a tarball named `team<num>.tar` using Autolab (`https://autolab.andrew.cmu.edu`). Untarring this file should give us a directory named `handin` which should contain the following:

Checkpoint 1:

- `Makefile` — Running `make` should produce an executable named `proxy`, as described in §2.2.3.

- `writeup.pdf` — This should contain the plots and analysis described above.

- `src` — A directory named `src` containing your source code. You may organize your code within this directory as you see fit.

3.3 Grading

The breakdown of grading for Checkpoint 1 is below.

Task	Weight	Subcriteria
Format	5%	*Assigned by human grader:* • Correct turnin – Makefile, compilation, properly tagged repo, TA's don't have to edit or search for files (5%)
Basic Proxy	30%	*Assigned based on tests:* • Receives requests from browser and forwards them to server. • Forwards the chunks it receives to the browser.
Bit Rate Adaptation	45%	*Assigned on tests:* • Implements throughput estimation • Bitrate selection
Writeup	20%	*Assigned by human grader:* • Correctness of graphs. • Discussion of the graphs.

4 Checkpoint 2: DNS Load Balancing

You MUST complete the following:

(1) Proxy. Implement a proxy that forwards requests received from the browser to the server, and forwards to chunk it receives to the browser. See §2.1 for details.

(2) Bitrate adaptation. Implement bit rate adaptation as descibed in Section 2.2.

(3) Explore the behavior of your two DNS redirect versions. Using the proxy and your DNS server, run the same set of tests you ran for CP1. Do the tests both for the RR version of the DNS server and the version that uses OSPF-based network distance. Discuss the differences you obvserver between the different runs, both with respect to the different values of α and the different DNS redirect algorithms.

4.1 Hand In

You will submit your code as a tarball named team<num>.tar using Autolab (https://autolab.andrew.cmu.edu). Untarring this file should give us a directory named handin which should contain the following:

- Makefile — Running make should produce an executable named proxy, as described in §2.2.3, and an executable named nameserver, as described in §2.3.2.

- writeup.pdf — This should contain the plots and analysis described above.

- src — A directory named src containing your source code. You may organize your code within this directory as you see fit, with the exception that we should find mydns.c immediately inside src.

4.2 Grading

The breakdown of grading for Checkpoint 1 is below.

Task	Weight	Subcriteria
Format	15%	*Assigned by human grader:* • Correct turnin – Makefile, compilation, properly tagged repo, TA's don't have to edit or search for files (5%) • Code style (5%) • Code commenting (5%)
DNS redirect	50%	*Assigned based on tests:* • DNS protocol • DNS resolution library • Round robin load balancing • LSA processing and "geographic" load distribution
Writeup	35%	*Assigned by human grader:* • Correctness of graphs and discussion on impact of α • Correctness of graphs and discussion on comparison RR versus OSPF-based redirect metric

5 Development Environment

For this project, we are providing a virtual machine pre-configured with the software you will need. We strongly recommend that you do all development and testing in this VM; your code must compile and run correctly on this image as we will be using it for grading. This section describes the VM and the starter code it contains.

The VM avalable is available from `https://cmu.app.box.com/s/ega2bntaqrnnqqb438tr2x1s723y71bj`.

5.1 Virtual Box

The virtual machine disk (VMDK) we provide was created using VirtualBox, though you may be able to use it with other virtualization software. The detailed process to set up the virtual machine and the image will be provided on our course websit. VirtualBox is a free download for Windows, OSX, and Linux on `https://www.virtualbox.org/wiki/Download_Old_Builds_5_0_pre18`. We've already set up an admin account:

Username: proj3

Password: proj3

Please make sure your download the version 5.0.16

5.2 Starter Files

You will find the following files in `/home/proj3/bitrate-project-starter`. This directory is a git repository; as we find bugs in the starter code and commit fixes, you can get the update versions with a `git pull`.

`mydns.h` The interface for your DNS resolution library (§2.3.2).

`common` Common code used by our network simulation and LSA generation scripts.

`lsa`

`lsa/genlsa.py` Generates LSAs for a provided network topology. You may use this to generate LSA files beyond those we provide, if you like.

`netsim`

`netsim/netsim.py` This script controls the simulated network; see §5.3.

`netsim/tc_setup.py` This script adjusts link characteristics (BW and latency) in the simulated network. It is called by `netsim.py`; you do not need to interact with it directly.

`netsim/apache_setup.py` This file contains code used by `netsim.py` to start and stop Apache instances on the IP addresses in your `.servers` file; you do not need to interact with it directly.

`grapher.py` A script to produce plots of link utilization, fairness, and smoothness from log files. (See §2.2.3.)

`topos`

`topos/topo1`

`topos/topo1/topo1.clients` A list of IP addresses, one per line, for the proxies. (Used by `netsim.py` to create a fake network interface for each proxy.)

`topos/topo1/topo1.servers` A list of IP addresses, one per line, for the video servers. (Used by your DNS server and by `netsim.py` to create a fake interface for each server.)

`topos/topo1/topo1.dns` A single IP address for your DNS server. (Used by `netsim.py` to create a fake interface for your DNS server.)

`topos/topo1/topo1.links` A list of links in the simulated network. (Used by `genlsa.py`.)

`topos/topo1/topo1.bottlenecks` A list of bottleneck links to be used in `topo1.events`. (See §5.3.)

`topos/topo1/topo1.events` A list of changes in link characteristics (BW and latency) to "play." See the comments in the file. (Used by `netsim.py`.)

`topos/topo1/topo1.lsa` A list of LSAs heard by the DNS server in this topology.

`topos/topo1/topo1.pdf` A picture of the network.

`topos/topo2`

`...`

5.3 Network Simulation

To test your system, you will run everything (proxies, servers, and DNS server) on a simulated network in the VM. You control the simulated network with the `netsim.py` script. You need to provide the script with a directory containing a network topology, which consists of several files. We provide two sample toplogies; feel free to create your own. See §5.2 for a description of each of the files comprising a topology. Note that `netsim.py` requires that each constituent file's prefix match the name of the topology (e.g., in the `topo1` directory, files are named `topo1.clients`, `topo1.servers`, etc.).

To start the network from the `netsim` directory:

$$\text{./netsim.py <topology> start}$$

Starting the network creates a fake network interface for each IP address in the `.clients`, `.servers`, and `.dns` files; this allows your proxies, Apache instances, and DNS server to bind to these IP addresses.

To stop it once started (thereby removing the fake interfaces), run:

$$\text{./netsim.py <topology> stop}$$

To facilitate testing your adaptive bitrate selection, the simulator can vary the bandwidth and latency of any link designated as a bottleneck in your topology's `.bottlenecks` file. (Bottleneck links must be declared because our simulator limits you to adjusting the characteristics of only one link between any pair of endpoints. This also means that some topologies simply cannot be simulated by our simulator.) To do so, add link changes to the `.events` file you pass to `netsim.py`. Events can run automatically according to timings specified in the file or they can wait to run until triggered by the user (see `topos/topo1/topo1.events` for an example). When your `.events` file is ready, tell `netsim.py` to run it:

<p align="center"><code>./netsim.py <topology> run</code></p>

Note that you must start the network before running any events. You can issue the `run` commands as many times as you want without restarting the network. You may modify the `.events` file between runs, but you must *not* modify any other topology files, including the `.bottlenecks` file, without restarting the network. Also note that the links stay as the last event configured them even when `netsim.py` finishes running.

Link State Advertisements You will notice that the simulated network described above doesn't actually contain any routers — we've taken them out of the picture to make your lives simpler by allowing you to independently set the bandwidth between any pair of endpoints (fake NICs). But, for the sake of the LSAs your DNS server needs to process, we need to *pretend that there are fake routers*. Yes, this is a tad confusing and contrived. Sorry. Deal with it. So, in each sample topology, we provide a `.links` file containing pairs of network elements (hosts or routers) between which there is a link in this pretend fake network. If the endpoint is a host, use its fake IP address in the `.links` file. If it is a pretend fake router, we use any string (e.g., "router1").

If want to make a new LSA file based on a new network topology for your testing, you can generate one using the `genlsa.py` script (run `./genlsa.py -h` for information on how to use it).

5.4 Apache

You will use the Apache web server to server the video files. `netsim.py` automatically starts an instance of Apache for you on each IP address listed in your topology's `.servers` file. Each instance listens on port 8080 and is configured to serve files from `/var/www`; we have put sample video chunks here for you.