

# 15-441/641: Content Delivery and Peer-to-Peer

15-441 Fall 2019  
 Profs **Peter Steenkiste** & Justine Sherry


Fall 2019  
<https://computer-networks.github.io/fa19/>



**Carnegie  
 Mellon  
 University**

## Problem: Scaling Content Delivery

- Millions of clients  $\Rightarrow$  server and network meltdown




## Outline

- Peer-to-peer
  - Overlays: naming, addressing, and routing
- CDNs
- Load balancing – consistent hashin



## P2P System



- Leverage the resources of client machines (peers)
  - Computation, storage, bandwidth



## P2P Definition

*Distributed systems consisting of interconnected nodes able to **self-organize** into network topologies with the purpose of **sharing resources** such as content, CPU cycles, storage and bandwidth, **capable of adapting to failures** and accommodating transient populations of nodes while maintaining acceptable connectivity and performance, **without requiring the intermediation or support of a global centralized server or authority.***

– A Survey of Peer-To-Peer Content Distribution Technologies,  
Androutsellis-Theotokis and Spinellis



## Why p2p?

- Harness lots of spare capacity
  - 1 Big Fast Server: \$10k/month++ versus 1000s .. 1000000s clients: \$ ??
  - Capacity grows with the number of users!
- Build very large-scale, self-managing systems
  - Same techniques useful for companies,
    - E.g. Akamai's 14,000+ nodes, Google's 100,000+ nodes
    - But: servers vs. arbitrary nodes, hard vs. soft state (backups vs caches), ....
    - Also: security, fairness, freeloading, ..
- No single point of failure
  - Some nodes go down – others take over
  - ... government shuts down nodes – peers in other countries are available




## Key Idea: Network Overlay

- A network overlay is a network that is layered on top of the Internet
  - Simplified picture: overlays use IP as their datalink layer
- Overlays need the equivalent of all the functions IP networks need:
  - Naming and addressing
  - Routing
  - Bootstrapping
  - Security, error recovery, etc.



## P2P Construction



## Names, addresses, and routing

### The Internet


- Endpoint: host
- Name: hierarchical domain name
- Address: IP address of node that has the content, plus content name
- Routing: how to reach host, e.g., BGP, ...

### Content retrieval:

- End-point: content
- Name: identifies content you are looking for
- E.g., hash of file, key words
- Address: the IP address of node that has the content, plus content name
- Routing: how to find the data



## Common P2P Framework




## What is (was) out there?

	Central	Flood	Super-node flood	Route
Whole File	Napster	Gnutella		Freenet
Chunk Based	BitTorrent		KaZaA (bytes, not chunks)	DHTs eDonkey 2000



## Napster: Central Database



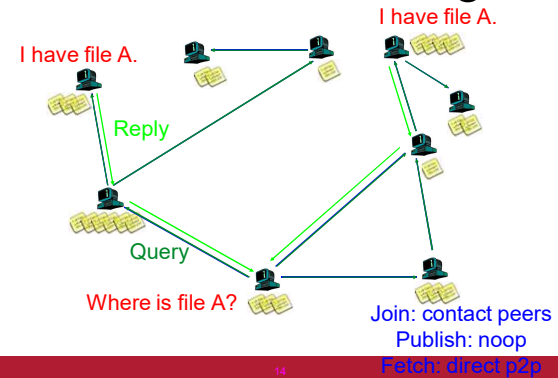
## Napster: Discussion

- Pros:
  - Simple
  - Search scope is  $O(1)$
  - Controllable (pro or con?)
- Cons:
  - Server maintains  $O(N)$  State
  - Server does all processing
  - Single point of failure



13

## Gnutella: Flooding



14

## Gnutella: Discussion

- Pros:
  - Fully de-centralized
  - Search cost distributed
  - Processing @ each node permits powerful search semantics
- Cons:
  - Search scope is  $O(N)$
  - Search time is  $O(???)$
  - Nodes leave often, network unstable
- TTL-limited search works well for haystacks.
  - For scalability, does NOT search every node.
  - May have to re-issue query later



15


## KaZaA: Query Flooding

- First released in 2001 and also very popular
- **Join:** on startup, client contacts a "supernode" ... may at some point become one itself
- **Publish:** send list of files to supernode
- **Search:** send query to supernode, supernodes flood query amongst themselves.
- **Fetch:** get the file directly from peer(s); can fetch simultaneously from multiple peers



16

## KaZaA: Intelligent Query Flooding



17

## KaZaA: Discussion

- Works better than Gnutella because of query consolidation
- Several nodes may have requested file... How to tell?
  - Must be able to distinguish identical files
  - Same filename not necessarily same file...
- Use Hash of file
  - Can fetch bytes [0..1000] from A, [1001...2000] from B
- Pros: Tries to take into account node heterogeneity:
  - Bandwidth, computational resources, ...
- Cons: Still no guarantees on search scope or time
- Challenge: want stable superpeers – good prediction
  - Must also be capable platforms



20


## BitTorrent: Swarming

- Started in 2001 to efficiently support flash crowds
  - Focus is on fetching, not searching
- **Publish:** Run a tracker server.
- **Search:** Find a tracker out-of-band for a file, e.g., Google
- **Join:** contact central "tracker" server for list of peers.
- **Fetch:** Download chunks of the file from your peers. Upload chunks you have to them.
- Comparison with earlier architectures:
  - Focus on fetching of "few large files"
  - Chunk based downloading
  - Anti-freeloading mechanisms




22

## BitTorrent: Publish/Join



23

## BitTorrent: Fetch



24

## BitTorrent: Summary

- Pros:
  - Works reasonably well in practice
  - Gives peers incentive to share resources; avoids freeloaders
- Cons:
  - Pareto Efficiency relative weak condition
  - Central tracker server needed to bootstrap swarm
  - (Tracker is a design choice, not a requirement, as you know from your projects. Could easily combine with other approaches.)



25

## When are p2p Useful?

- Works well for caching and “soft-state”, read-only data
  - Works well! BitTorrent, KaZaA, etc., all use peers as caches for hot data
- Difficult to extend to persistent data
  - Nodes come and go: need to create multiple copies for availability and replicate more as nodes leave
- Not appropriate for search engine styles searches
  - Complex intersection queries (“the” + “who”): billions of hits for each term alone
  - Sophisticated ranking: Must compare many results before returning a subset to user
  - Need massive compute power



27


## Outline

- Peer-to-peer
  - Overlays: naming, addressing, and routing
- CDNs
- Load balancing – consistent hashin




28

## Content Delivery: Possible Bottlenecks



## Reminder: Caching with Forward Proxies

- Cache documents close to **clients**  
→ decrease latency
- Typically done by ISPs or enterprises  
→ reduce provider traffic load
- CDNs proactively cache for the content providers (their clients)
- Typically cache at different levels in the Internet hierarchy:
  - Last mile ISPs for low latency
  - Closer to core for broader coverage




## Content Distribution Networks (CDNs)

- The content providers are the CDN customers.

### Content replication

- CDN company installs hundreds of CDN servers throughout Internet
  - Close to users
- CDN replicates its customers' content in CDN servers. When provider updates content, CDN updates servers



## What is the CDN?

- Edge Caches: work with ISP and networks everywhere to install edge caches
  - Edge = close to customers
- Content delivery: getting content to the edge caches
  - Content can be objects, video, or entire web sites
- Mapping: find the "closest" edge server for each user and deliver content from that server
  - Network proximity not the same as geographic proximity
  - Focus is on performance as observed by user (quality)

## Potential Benefits

- Very good scalability
  - Near infinite if deployed properly
- Good economies at large scales
  - Infrastructure is shared efficiently by customers
  - Statistical multiplexing: hot sites use more resources
- Can reduce latency – more predictable performance
  - Through mapping to closest server
  - Avoids congestion and long latencies
- Can be extremely reliable
  - Very high degree of redundancy
  - Can mitigate some DoS attacks



33

## Server Selection

- Which server?
  - Lowest load: to balance load on servers
  - Best performance: to improve client performance
    - Based on Geography? RTT? Throughput? Load?
  - Any alive node: to provide fault tolerance
- How to direct clients to a particular server?
  - As part of naming: DNS redirect
  - As part of application: HTTP redirect
  - As part of routing: anycast, cluster load balancing



34

## Finding the “Closest Edge Cache – Example: Akamai DNS Redirect

- Akamai creates new domain names for each client
  - e.g., [a128.g.akamai.net](http://a128.g.akamai.net) for [cnn.com](http://cnn.com)
- The CDN's DNS servers are authoritative for the new domains
- The client content provider modifies its embedded URLs (= names) to reference the new domains – “Akamaize” content
  - e.g.: <http://www.cnn.com/image-of-the-day.gif> becomes
  - [http:// a128.g.akamai.net/image-of-the-day.gif](http://a128.g.akamai.net/image-of-the-day.gif) – name in the overlay
- **Requests now sent to CDN's infrastructure...**
- Generates and address: IP address of server + URI (tuple)
- Routing inside Akamai system identifies right replica to route to
  - IP takes care of rest once a replica has been selected (overlay!)



Effectively another layer of routing:  
the path your connection takes is  
redirected using DNS.






## Alternative Approaches

- Routing based (IP anycast)
  - Multiple CDN instances advertise the same IP address block
  - BGP will route packets to the closest one (fewest AS hops)
  - **Pros:** Transparent to clients, works when browsers cache failed addresses, circumvents many routing issues
  - **Cons:** Little control, complex, scalability, TCP can't recover
- Application based (HTTP redirects)
  - Send request to origin HTTP server which redirects the HTTP request to a CDN instance closer to the client
  - **Pros:** Application-level, fine-grained control
  - **Cons:** Additional load and RTTs, hard to cache, availability concerns




## Process Flow



1. User wants to download distributed web content




## Process Flow



2. User is directed through Akamai's dynamic mapping to the "closest" edge cache




## Process Flow



3. Edge cache searches local hard drive for content




## Process Flow



3b. If requested object is not on local hard drive, edge cache checks other edge caches in same region for object




## Process Flow



3b. If requested object is not cached or not fresh, edge cache sends an HTTP GET the origin server




## Process Flow



3c. Origin server delivers object to edge cache over optimized connection




## Process Flow



4. Edge server delivers content to end user




## Core Hierarchy Regions



1. User requests content and is mapped to optimal edge Akamai server




## Core Hierarchy Regions



2. If content is not present in the region, it is requested from most optimal core region




## Core Hierarchy Regions



3. Core region makes one request back to origin server



## Core Hierarchy Regions



4. Core region can serve many edge regions with one request to origin server



## Core CDN Features

Reduces traffic back to origin server

- Reduces infrastructure needs of customer
- Provides best protection against flash crowds
  - Especially important for large files (e.g. Operating System updates or video files)

Improved end-user response time

- Core regions are well connected
- Optimized connection speeds object delivery



## Outline

- Peer-to-peer
  - Overlays: naming, addressing, and routing
- CDNs
- Load balancing – consistent hashin



52

## Distributing Load across Servers

- Given document XYZ, we need to choose a server to use
  - E.g., in a data center
- Suppose we use simple hashing: modulo  $n$  of a hash of the name of the document
- Number servers from  $1 \dots n$ 
  - Place document XYZ on server  $(XYZ \bmod n)$
  - What happens when a servers fails?  $n \rightarrow n-1$ 
    - Same if different people have different measures of  $n$
  - Why might this be bad?



53

## Consistent Hash: Goals


- “view” = subset of all hash buckets that are candidate locations
  - Correspond to a real server
- Desired features
  - Load – all hash buckets have a similar number of objects assigned to them
  - Smoothness – little impact on hash bucket contents when buckets are added/removed
  - Spread – small set of hash buckets that may hold an object regardless of views



54

## Consistent Hash – Example


- Construction
  - Assign each of  $C$  hash buckets to random points on mod  $2^n$  circle, where, hash key size =  $n$ .
  - Map object to random position on unit interval
  - Hash of object = closest bucket
- Monotone → addition of bucket does not cause movement between existing buckets
- Spread & Load → small set of buckets that lie near object
- Balance → no bucket is responsible for large number of objects



55

## Consistent Hashing: Ring


- Use consistent has to map both **keys** and **nodes** to an  $m$ -bit identifier in the same (metric) identifier space
  - For example, use SHA-1 hashes
- Node identifier:** SHA-1 hash of IP address  
 IP="198.10.10.1" → SHA-1 → ID=123
- Key identifier:** SHA-1 hash of key  
 Key="LetItBe" → SHA-1 → ID=60
- Also need "rule" for assigning keys to nodes
  - For example: "closest", higher, lower, ..



56

## Consistent Hashing Example

**Rule:** A key is stored at its **successor**: node with next higher or equal ID



57

## Consistent Hashing Properties

- Load balance:** all nodes receive roughly the same number of keys
  - For  $N$  nodes and  $K$  keys, with high probability
    - Each node holds at most  $(1+\epsilon)K/N$  keys
    - Provided that  $K$  is large compared to  $N$
- When server is added, it receives its initial work load from "neighbors" on the ring
  - "Local" operation: no other servers are affected
  - Similar property when a server is removed



58

## Finer Grain Load Balancing

- Redirector knows all server IDs  $s_i$
- It can also track approximate “load” for more precise load balancing
  - Need to define load and be able to track it
- To balance load:
  - $W_i = \text{Hash}(\text{URL}, \text{ip of } s_i)$  for all  $i$
  - Sort  $W_i$  from high to low
  - Find first server with low enough load
- Benefits and drawbacks?



59

## Consistent Hashing Used in Many Contexts

- Distribute load across servers in a data center
  - The redirector sits in data center
- Finding storage cluster for an object in a CDN uses centralized knowledge
  - Why?
  - Can use consistent hashing in the cluster
- Consistent hashing can also be used in a distributed setting
  - P2P systems can use it find files (DHTs)



60